



# ONNX Introduction

---

Chin Huang, Cognitive OpenTech

*Data and AI  
Open Source Dojo*

# Outline

---

- ❖ Why and What is ONNX
- ❖ Overview and components
- ❖ ONNX ecosystem
- ❖ ONNX use cases

# Challenges and Issues

---

AI is booming!

- AI stacks evolved organically across organizations using a wide variety of frameworks
- Lots of target training and inferencing platforms
- Different hardware fabrics also evolved looking to meet training and serving needs

## Challenges

- How to translate from one framework/platform combination to another?
- Difficult to share and re-use domain knowledge and applications
- Increasing fragmentation across companies and teams
- Long AI development cycle, idea -> research -> dev -> production -> maintain -> refine -> upgrade

# Open and Interoperable AI



# What is ONNX (Open Neural Network eXchange)?

---

- Open ecosystem for **interchangeable AI models**
- Designed to be an **open specification and format**, empowering developers to freely select the framework/tool that works best for their project, at any stage of development.
- Commonly describes the model graph, which serves as an **Intermediate Representation (IR)** that captures the specific intent of the developer's source code.
- Enables flexible graph optimization.
- Persisted as **binary protobuf** files. Each model file contains the network structure and parameters of the model.
- Advantages: **open, framework neutral, cross-platform, complete ecosystem**

## Background and factoids

---

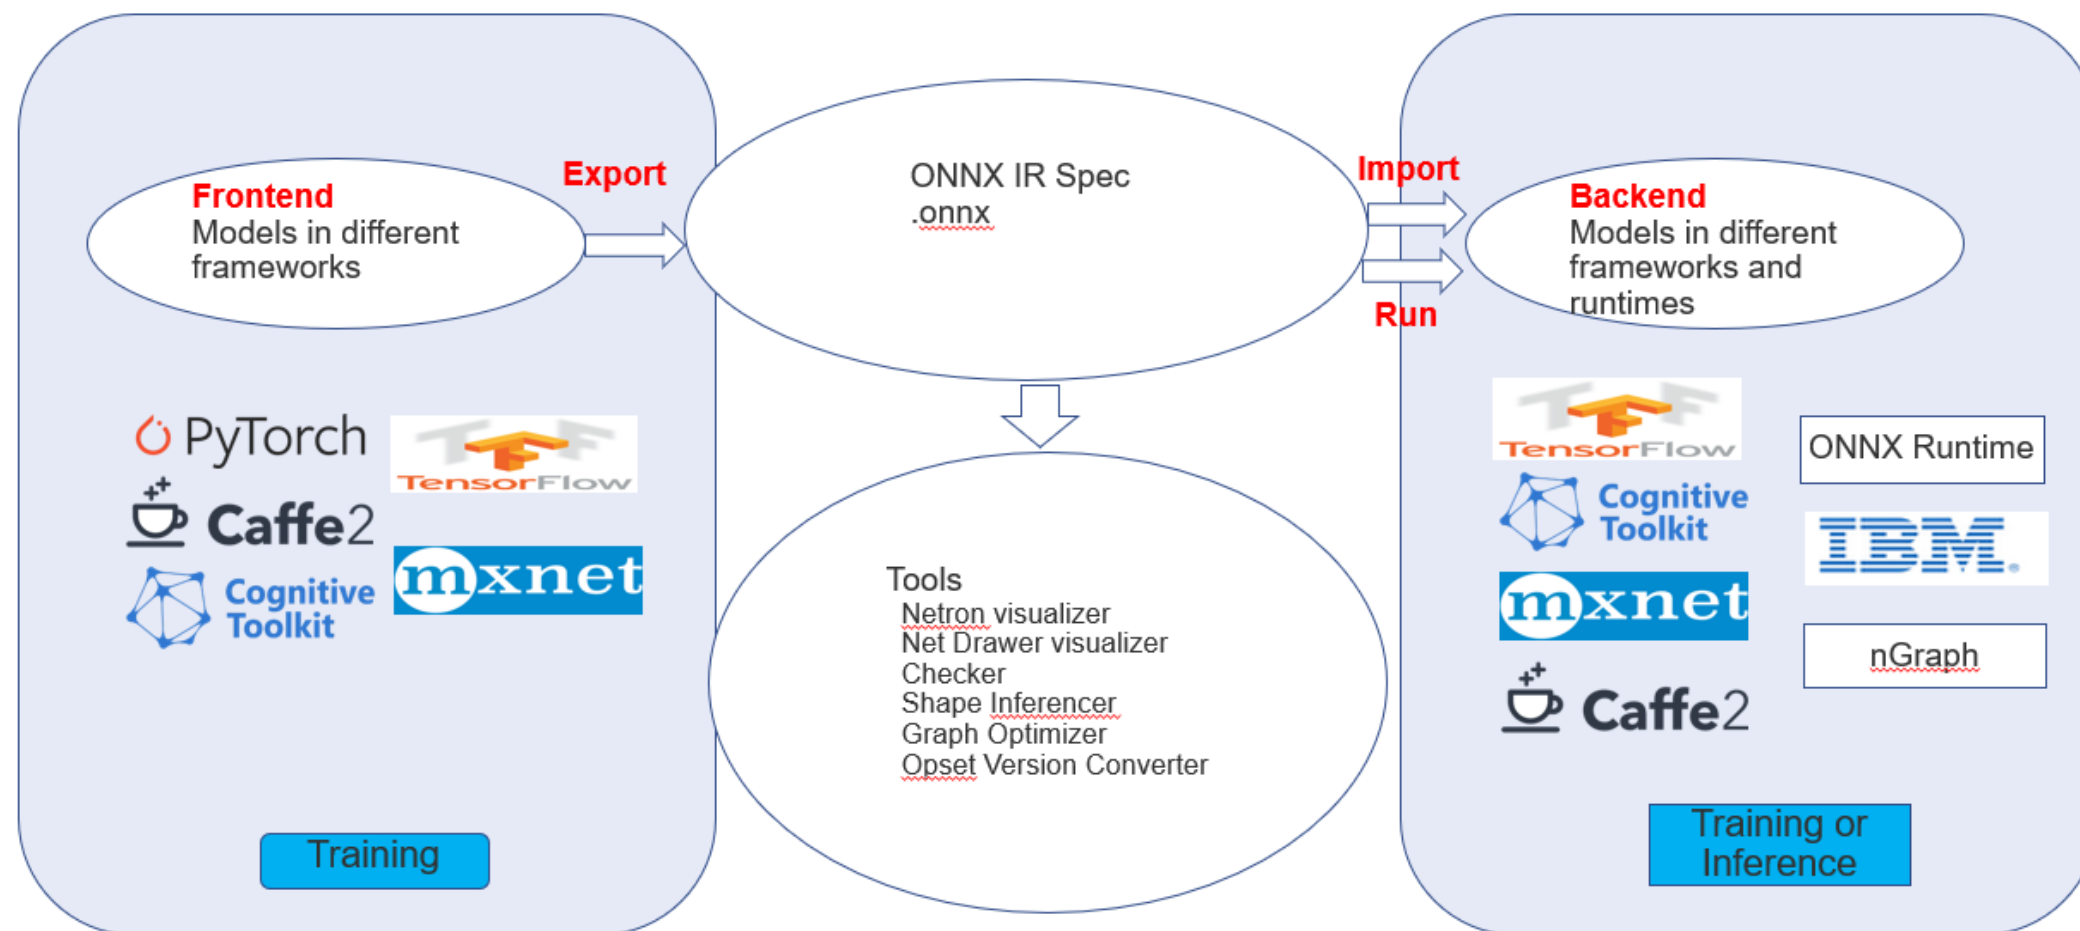
- Started Sept 2017 by Microsoft and Facebook
- Initial goal is to make it easier for data analysts to **exchange trained models** between different machine learning frameworks.
- Model training is added in 2020 to describe **trainable models**
- ONNX github has 20 repos. **onnx** is the core. Others are **tutorials, model zoo, importers and exporters** for frameworks.
- Onnx/onnx currently has 15 releases, 153 contributors, 8015 stars.
- Core is in C++ with python API and tools.
- Supported frameworks: Caffe2, Chainer, Cognitive Toolkit (CNTK), Core ML, MXNet, PyTorch, PaddlePaddle, Tensorflow...

# ONNX Partners



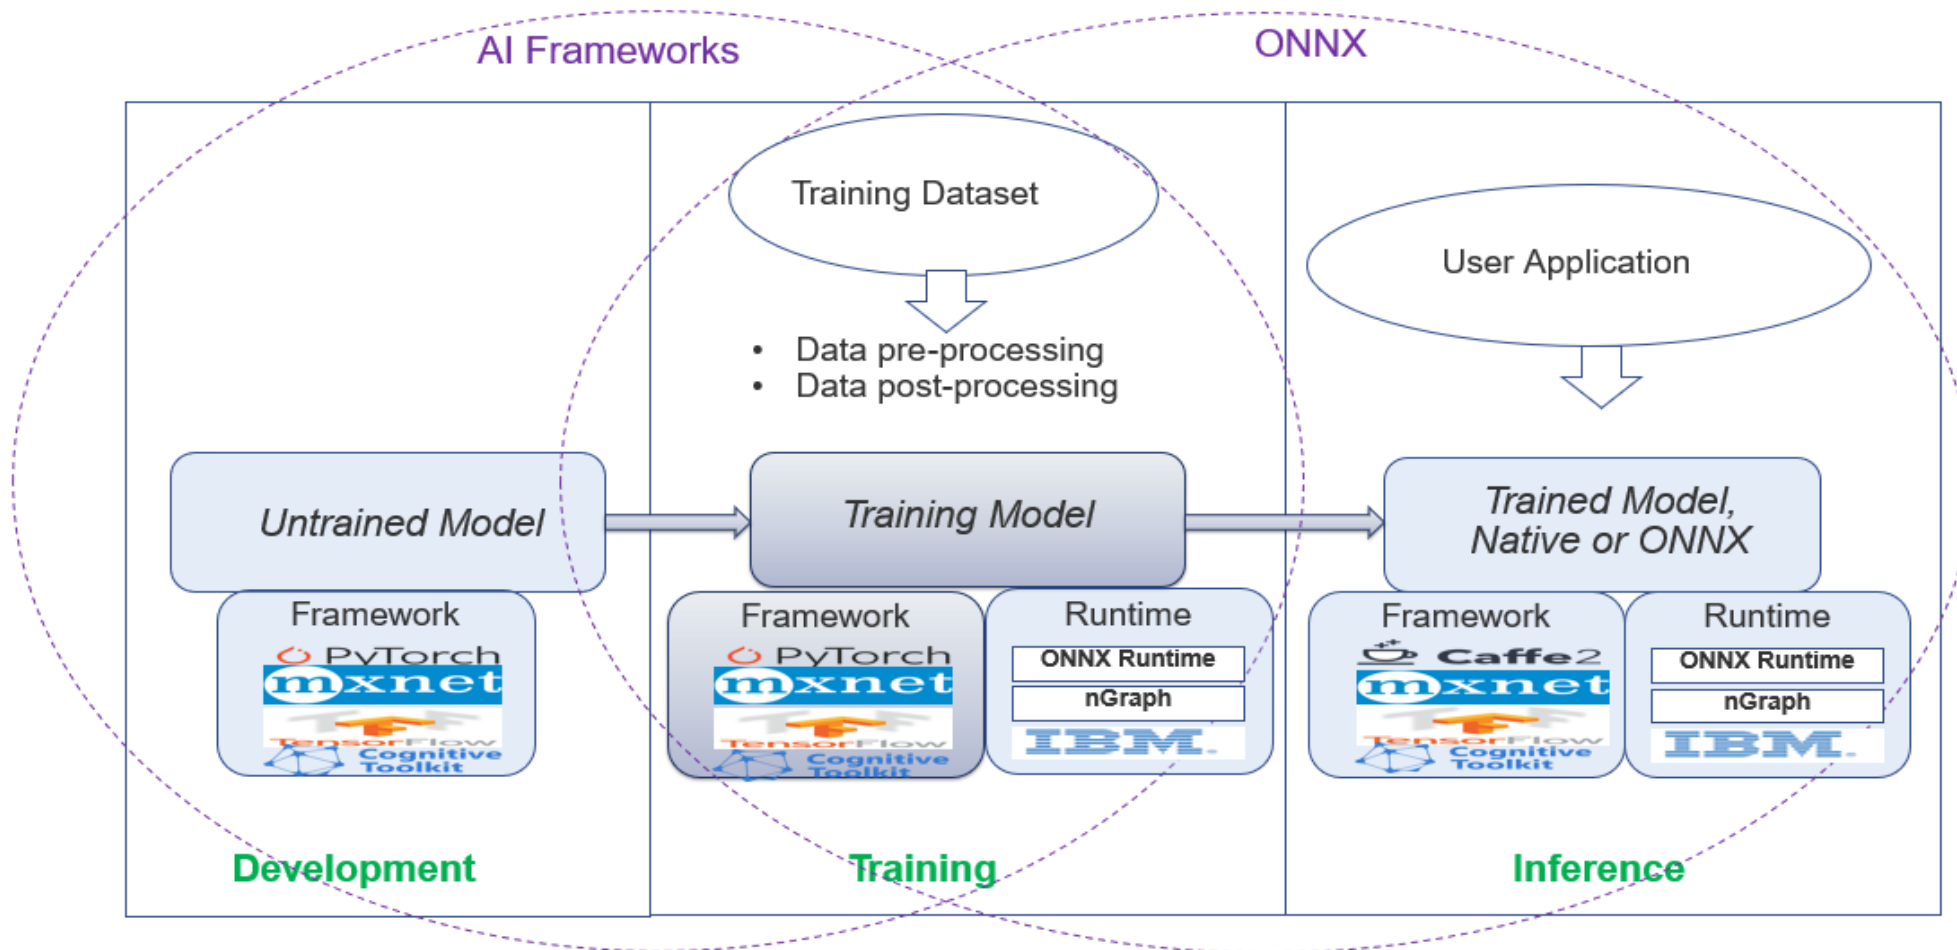


# ONNX Overview





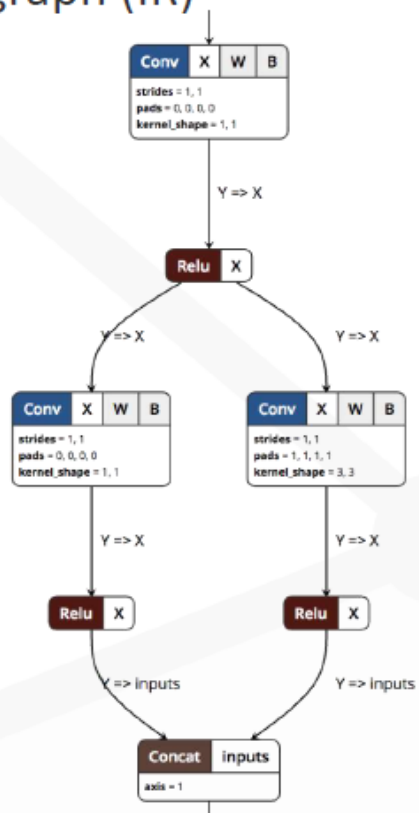
# AI Model Development Cycle with ONNX




# ONNX Core

ONNX is an open specification:

- Computation graph (IR)
- Data types
- Operators and functions



Two profiles:



# High Level IR

## Model

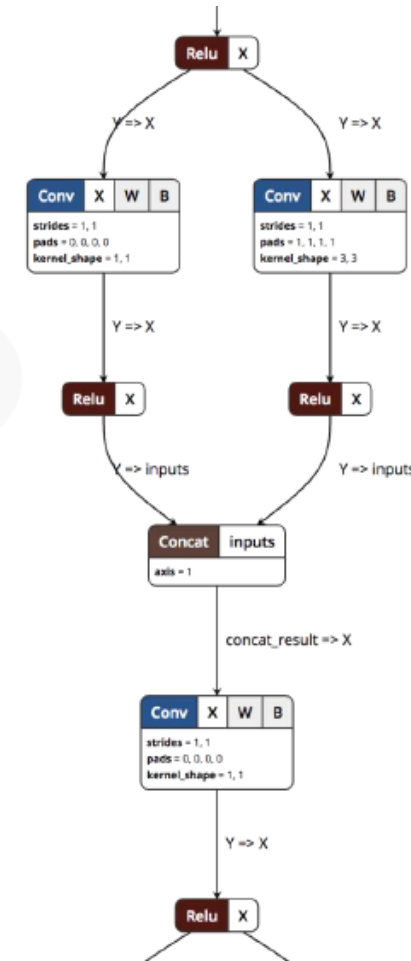
- Version info
- Metadata
- Acyclic computation dataflow graph

## Graph

- Inputs and outputs
- List of computation nodes
- Graph name

## Computation Node

- Zero or more inputs of defined types
- One or more outputs of defined types
- Operator
- Operator parameters



# Data Types

- Tensor type
  - Element types supported:
    - int8, int16, int32, int64
    - uint8, uint16, uint32, uint64
    - float16, float, double
    - bool
    - string
    - complex64, complex128
    - bfloat16
- Non-tensor types in ONNX-ML:
  - Sequence
  - Map

```
message TypeProto {
  message Tensor {
    optional TensorProto.DataType elem_type = 1;
    optional TensorShapeProto shape = 2;
  }
  // repeated T
  message Sequence {
    optional TypeProto elem_type = 1;
  };
  // map<K,V>
  message Map {
    optional TensorProto.DataType key_type = 1;
    optional TypeProto value_type = 2;
  };

  oneof value {
    Tensor tensor_type = 1;
    Sequence sequence_type = 4;
    Map map_type = 5;
  }
}
```

# Operators

An operator is identified by <name, domain, version>

## Core ops (ONNX, ONNX-ML, ONNX-Training domains)

- Should be supported by all ONNX-compatible products
- Generally cannot be meaningfully further decomposed
- Version incremented for breaking changes
- ONNX-Training ops for model training only

## Custom ops

- Ops specific to framework or runtime
- Indicated by a custom domain name
- Primarily meant to be a safety-valve

### Relu

Relu takes one input data (Tensor) and produces one output data (Tensor) where the rectified linear function,  $y = \max(0, x)$ , is applied to the tensor elementwise.

#### Version

This version of the operator has been available since version 6 of the default ONNX operator set. Other versions of this operator: Relu-1

#### Inputs

**x** : T  
Input tensor

#### Outputs

**y** : T  
Output tensor

#### Type Constraints

**T** : *tensor(float16), tensor(float), tensor(double)*  
Constrain input and output types to float tensors.

#### Examples


##### ▼ relu

```
node = onnx.helper.make_node(
    'Relu',
    inputs=['x'],
    outputs=['y'],
)
x = np.random.randn(3, 4, 5).astype(np.float32)
y = np.clip(x, 0, np.inf)

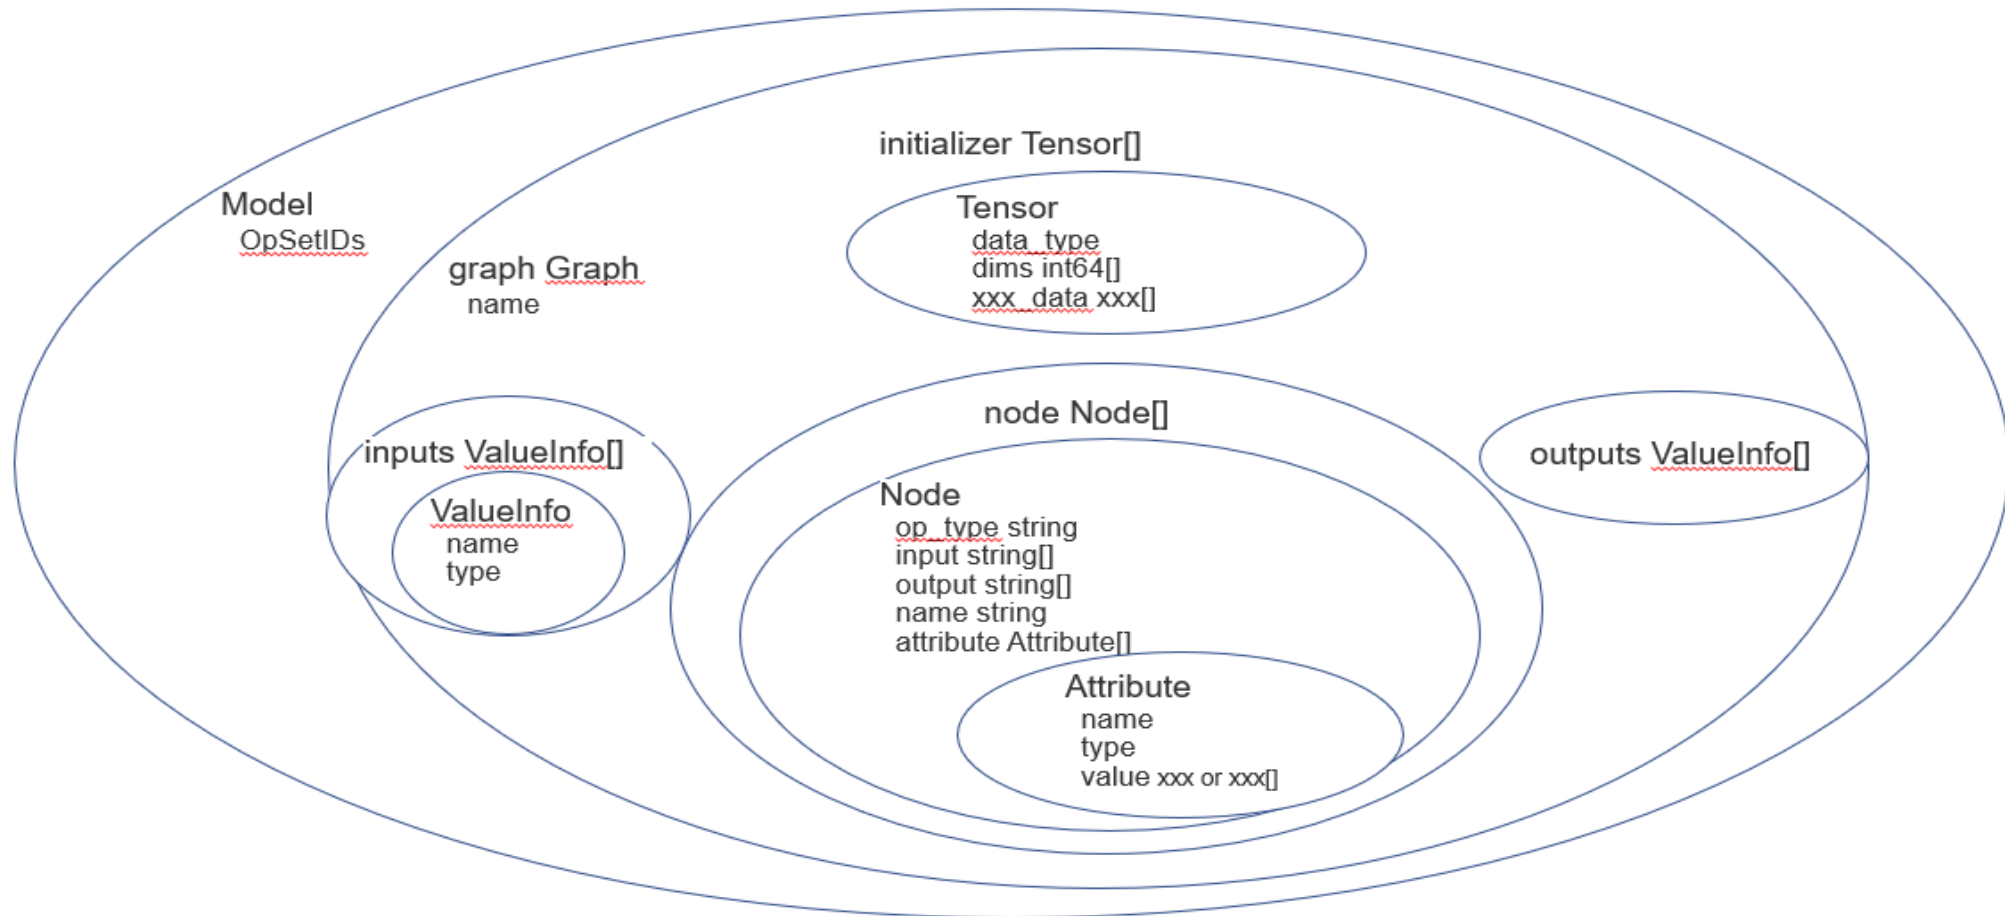
expect(node, inputs=[x], outputs=[y],
        name='test_relu')
```

# Functions

- Compound ops built with existing primitive ops
- Enable potential optimizations by runtimes, e.g. node fusion
- Two types of functions:
  - Standardized functions - defined by ONNX spec
  - In-model customized functions (not in spec yet)  
Defined within the model itself;  
runtimes/frameworks/tools that don't know the function can fallback to use the primitive ops

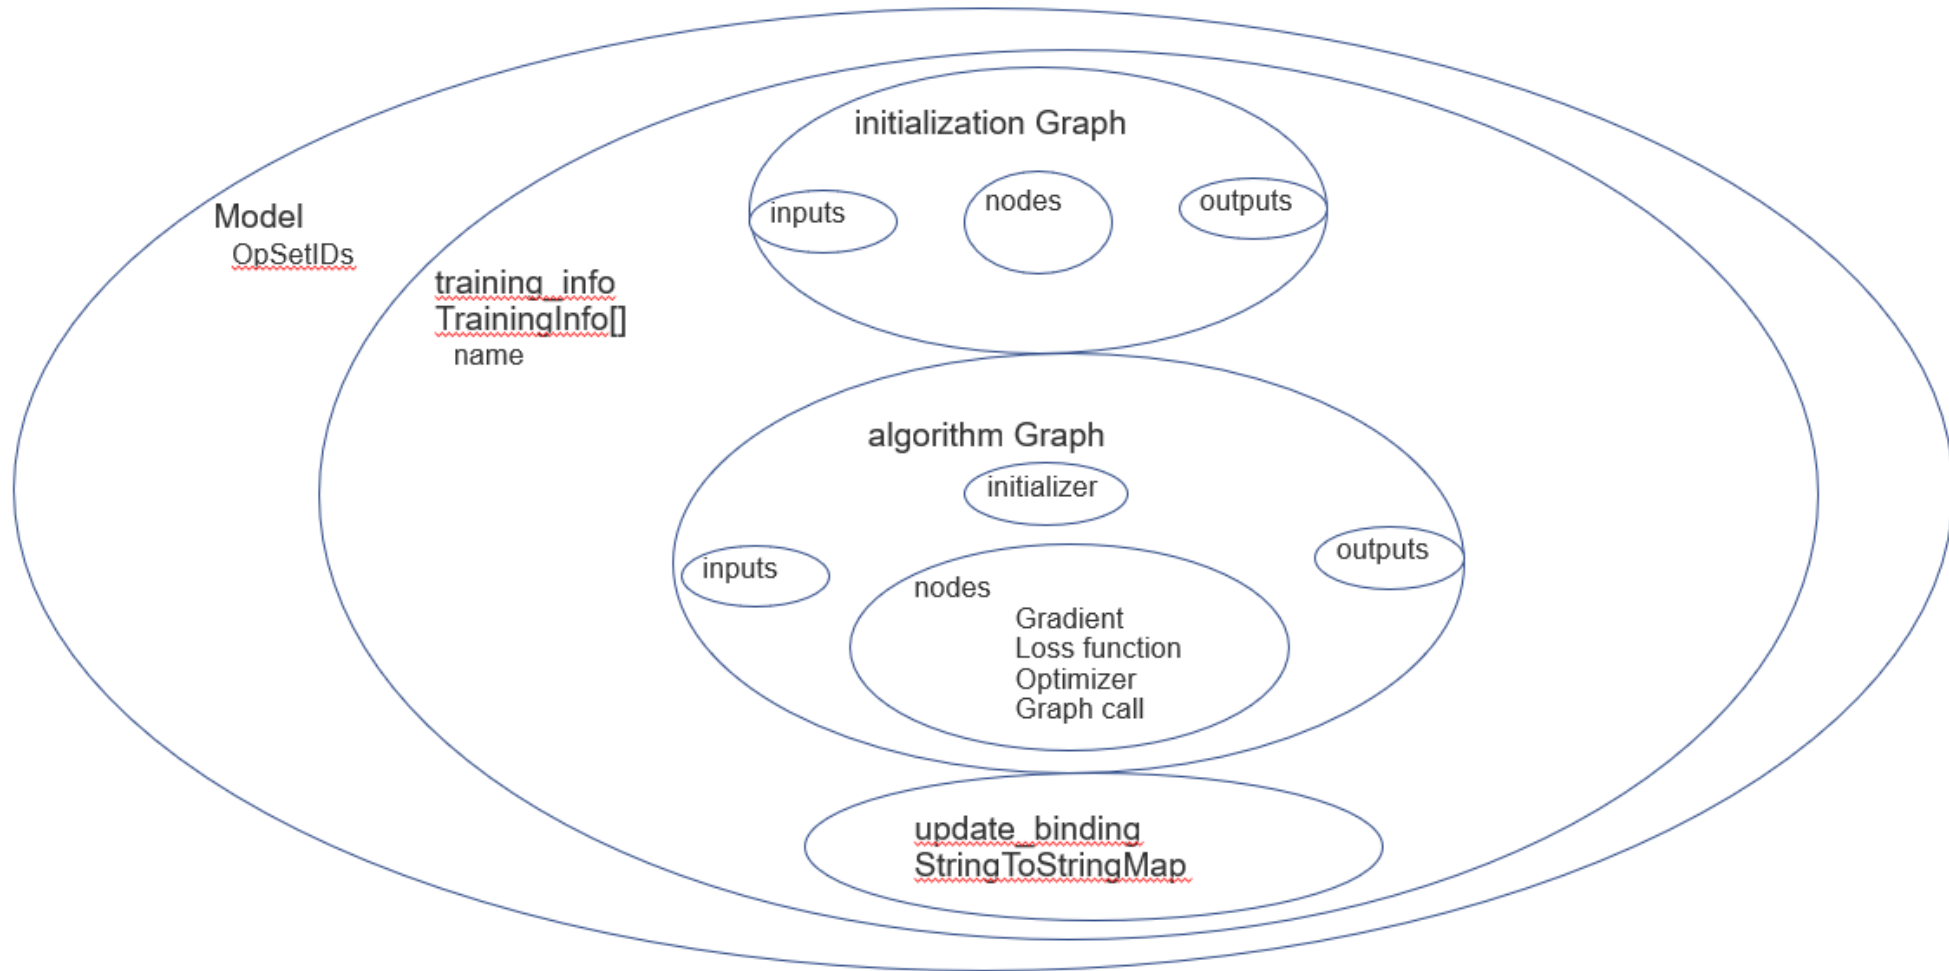


# ONNX IR Spec – Computation Graph

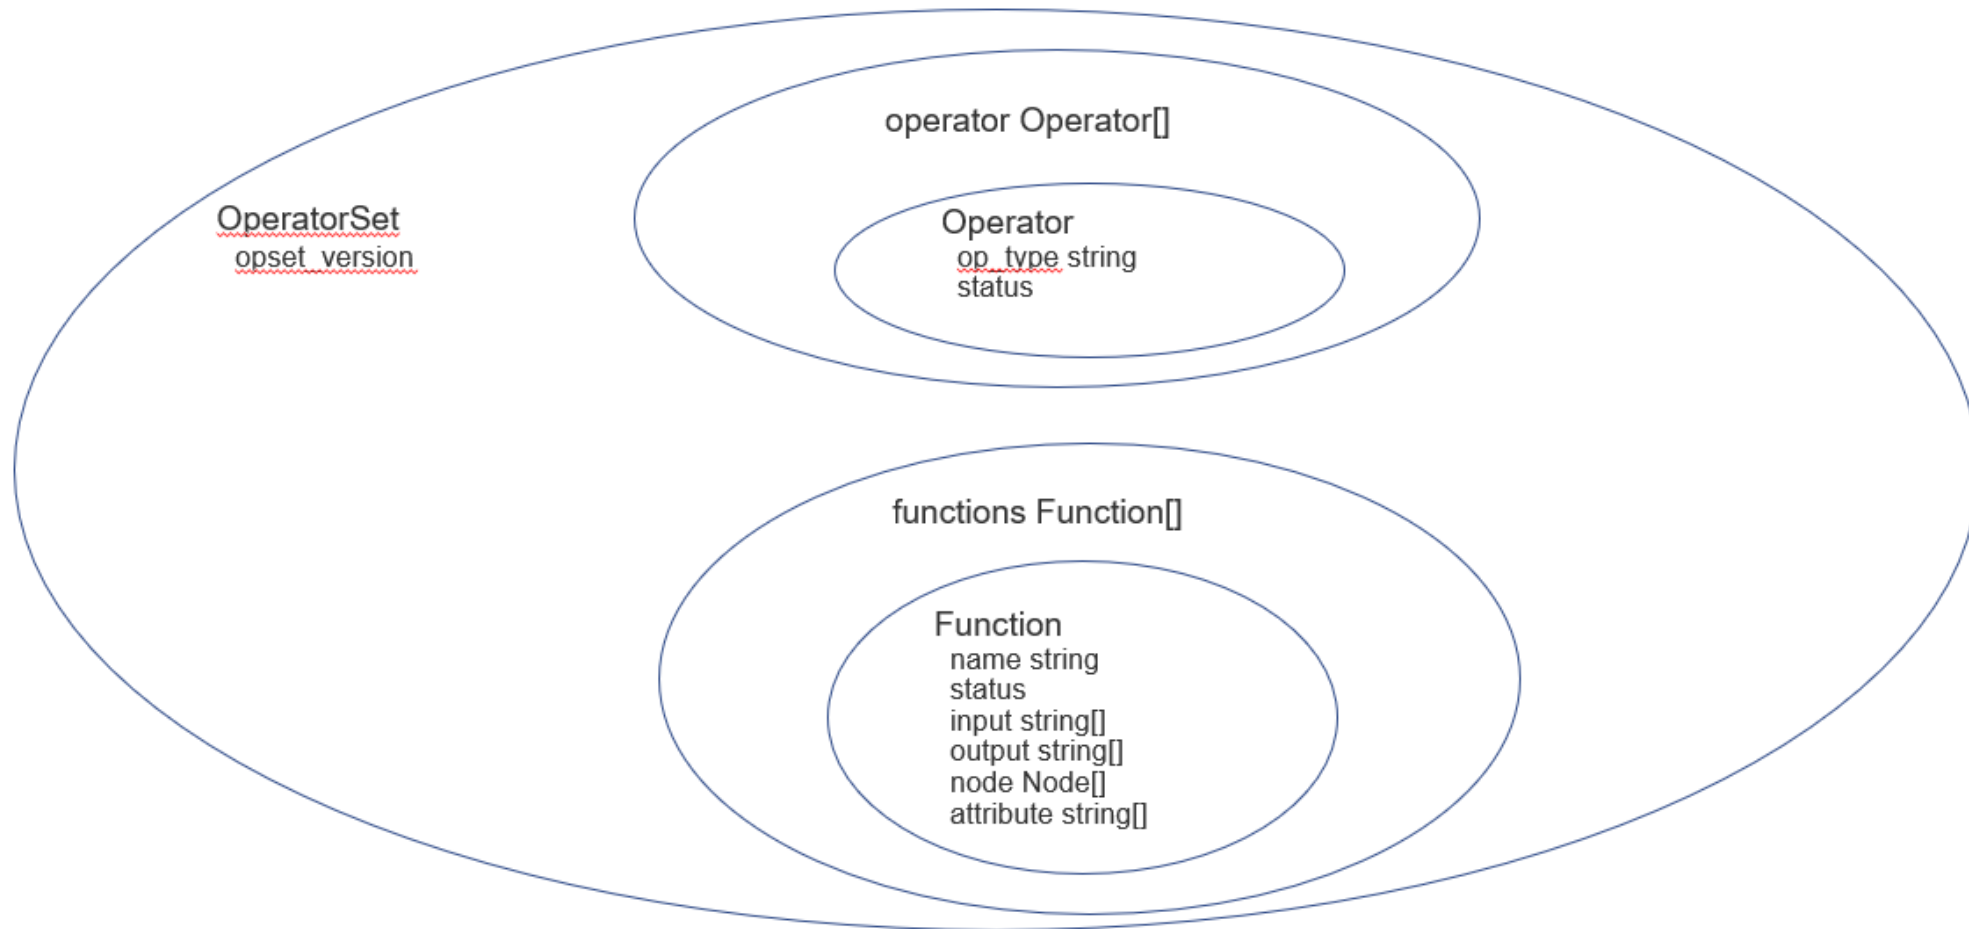




# ONNX IR Spec – Training Graph



# ONNX IR Spec – Operator Set



# ONNX Backend

---

- An ONNX backend is a library that can **run ONNX models**.
- Each AI framework can create a converter that **converts ONNX models** from and to the corresponding framework specific representation and then delegate the execution to the framework.
- ONNX has defined a unified (Python) backend interface at <https://github.com/onnx/onnx/blob/master/onnx/backend/base.py>
- ONNX provides a **standard backend test suite** to assist backend implementation verification. Two types of test are 1. nodes tests: verify whether a backend is performing the correct computation 2. model tests, verify the backend at the model level. It's strongly encouraged that each ONNX backend runs this test.

# ONNX Tools

---

- Visualizer
  - Net Drawer
  - Netron (<https://github.com/lutzroeder/netron>)
- Checker: validate ONNX models
- Optimizer: perform optimizations on ONNX models
- Shape inferencer: query the shape of tensor
- Version converter: convert ONNX models between different opset versions

# ONNX Optimizer

---

- ONNX provides a C++ library for performing arbitrary optimizations on ONNX models, as well as a growing list of [prepackaged optimization passes](#).
- The primary motivation is to share work between the many ONNX backend implementations.
- The aim is to provide all such passes along with ONNX so that they can be re-used with a single function call.
- The API Optimize() accepts an input ModelProto and a list of optimization passes to apply, and returns a new ModelProto which is the result of applying each of the named passes in sequence to the model.

# ONNX Interface for Framework Integration (ONNXIFI)

---

- A **cross-platform API** for loading and executing ONNX graphs on optimized backends
- Features:
  - **Standardized interface** for neural network inference on special-purpose accelerators, CPUs, GPUs, DSPs, and FPGAs
  - Dynamic discovery of available backends for model execution
  - Dynamic discovery of supported ONNX Operators on each backend
  - **Extensible interface** for vendor-specific operators and vendor-provided ONNXIFI implementation

# Python API

---

- ONNX model lifecycle support
  - Loading, saving, creating models
  - Checking/validating
  - Optimizing
  - Running shape inference
  - Converting versions
  - Polishing (runs model checker, optimizer, shape inference engine on the model, and also strips the doc\_string)



# ONNX Use Cases

- Microsoft

## PLATFORMS



AzureML



WinML



ML.Net

## PRODUCTS



Bing



Microsoft  
Cognitive Services

Office 365



ads



Power BI



Up to  
14.6x

Performance gains seen by  
Microsoft services

100s of  
Millions

# of devices where ONNX  
Runtime is running

Billions

# of requests handled by ONNX  
Runtime across Microsoft  
services

# ONNX Use Cases

---

- Huawei
  - MindSpore: a full-stack AI framework open sourced in 1Q 2020
  - MindSpore IR -> ONNX model -> ONNX Runtime
  - ONNX model -> MindSpore IR, optimized for Ascent chips
- MathWorks (MATLAB)
  - Two-way model exchange with ONNX <--> MATLAB
  - Automatic code gen, visualization, re-training in MATLAB
  - Goals: Import 90% of ONNX model zoo models to MATLAB, export 90% of MATLAB models to ONNX!

# ONNX Use Cases

---

- Intel OpenVINO Toolkit
  - Visual Inference & Neural Network Optimization, computer vision toolkit for edge computing
  - Use case 1: direct convert from ONNX models
  - Use case 2: as an ONNX Runtime EP (execution provider)
- Tencent (Chinese social platforms such as QQ and WeChat/Weixin)
  - <https://github.com/Tencent/ncnn>
  - NCNN deploys onnx for mobile, supporting Arm and GPUs
  - Sample case: pytorch -> onnx -> ncnn

# ONNX Use Cases

---

- Alibaba
  - SinianAI: a full-stack AI framework, acceleration for Cloud, Edge, and IoT
  - Use cases: smoking detection, self checkout, voice activation
- UC Santa Cruz
  - Deep neural network using ONNX for efficient genome analysis
  - Global effort with multiple cross-country research consortiums

# ONNX at IBM

---

- Open Source Contributions
  - ONNX
    - Converters SIG lead
    - Training working group lead
    - Power CI build and test
    - Release manager for 1.7
  - ONNX-TensorFlow converter
  - ONNX-MLIR lead
- Open Governance: Took ONNX to LF AI as a graduated project!

# ONNX at IBM

---

- IBM Products and Services
  - IBM Watson Machine Learning for Z
  - IBM Watson Machine Learning Accelerator